

GEONHWA JEONG

Klaus Advanced Computing Building 3305, 266 Ferst Dr NW, Atlanta, GA 30332, USA.

☎ +1 470-309-8607 ✉ geonhwa.jeong@gatech.edu 🌐 ghjeong12.github.io

RESEARCH INTERESTS

Computer architecture, HW/SW co-design, Domain-specific accelerators, Compiler optimization, Reconfigurable architecture, Datacenter-scale computing, Deep learning

EDUCATION

Ph.D. Student in Computer Science

Aug. 2019 - Current

Advisor: Prof. Tushar Krishna

GPA: 4.0/4.0

Georgia Institute of Technology (Georgia Tech)

Master of Science in Computer Science

Aug. 2019 - Dec. 2021

Specialization: Machine Learning

GPA: 4.0/4.0

Georgia Institute of Technology (Georgia Tech)

Bachelor of Science in Creative IT Engineering

Mar. 2013 - Feb. 2019

Double Major in Computer Science and Engineering

Summa Cum Laude

Pohang University of Science and Technology (POSTECH)

RESEARCH EXPERIENCES

Synergy Lab at Georgia Tech

Nov. 2019 - Current

Research Assistant (Advisor: Prof. Tushar Krishna)

Atlanta, USA

To fully exploit available parallelism in spatial accelerators, it is important to schedule and map the input data efficiently to the processing elements maximizing data reuse. I worked on optimizing the dataflow considering both software and hardware sides [*PACT'21, TPDS'21, MICRO'20*]. Also, I led a project to efficiently integrate a systolic array based matrix engine in CPU to accelerate deep learning workloads [*DAC'21*]. Currently, I am leading a project to extend the previous work to leverage sparsity in DNNs.

High Performance Architecture Lab at Georgia Tech

Sep. 2019 - Nov. 2019

Research Student (Advisor: Prof. Hyesoon Kim)

Atlanta, USA

The Unified Virtual Memory with GPU and CPU not only gets rid of programmers burden a lot but also enables running a program with the working set size larger than the GPU capacity. I worked on efficient page prefetching mechanisms with various workloads.

Compiler Optimization Research Lab at POSTECH

Sep. 2017 - June 2018

Research Student (Advisor: Prof. Hanjun Kim)

Pohang, Republic of Korea

I was engaged in three projects including implementation of IoT platform for various types of users (manufacturer, service developer, user), development of programmable magnetic blocks to teach computational thinking to kids and development of hot function/basic block detector using LLVM compiler.

Database and Data Mining Lab at POSTECH*Research Student (Advisor: Prof. Wook-Shin Han)*

Mar. 2015 - Nov. 2015

Pohang, Republic of Korea

I was involved in a team to develop a new graph database system which can process a large amount of streaming data. I participated in the initial design of the query processing engine and the development of memory management system to support the engine [SIGMOD'18].

WORK EXPERIENCES

NVIDIA Research*Research Intern*

May 2022 - Aug 2022

Santa Clara, USA

I worked on a project to accelerate DNN model inferences using tensor approximation while minimizing accuracy drop.

Facebook*Research Engineering Intern*

May 2021 - Aug 2021

Seattle, USA

I worked on workload characterization for data compression in datacenter-scale services at Facebook and explored HW offloading opportunities [ISPASS'22].

Intel Labs*Graduate Technical Intern*

May 2020 - Aug 2020

Oregon, USA

I implemented a performance model for a new architectural feature of CPU and built a framework to validate the model by comparing against RTL results.

VoyagerX*Software Engineering Intern*

Feb. 2019 - July 2019

Seoul, Republic of Korea

I developed a mobile application for automatic meeting notes with speaker diarization using a deep learning model to extract feature vectors from voice data.

Samsung Research*Field Training*

July 2018 - Aug. 2018

Seoul, Republic of Korea

I was involved in a team developing a cloud management system, which is used by the members of Samsung Research to develop their own programs, in GitOps manner to keep the more robust system.

Korea Augmentation to the US Army*Service Member*

Nov. 2015 - Aug. 2017

Seoul, Republic of Korea

I served mandatory military service as a movement specialist with the United States Army and participated in three combined exercises (KR 17, UFG 17, and KR 18).

Kakao Corp.*Intern*

Jan. 2015 - Feb. 2015

Seongnam, Republic of Korea

I developed an abuse detection system to automatically identify abusers to prevent general users from being exposed to inappropriate posts and comments.

PUBLICATIONS

- [1] **Geonhwa Jeong**, Bikash Sharma, Nick Terrell, Abhishek Dhanotia, Zhiwei Zhao, Niket Agarwal, Arun Kejariwal, and Tushar Krishna, “Understanding Data Compression in Warehouse-Scale Datacenter Services,” in *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, May 2022.
- [2] Gordon E. Moon, Hyoukjun Kwon, **Geonhwa Jeong**, Prasanth Chatarasi, Sivasankaran Rajamanickam, Tushar Krishna, “Evaluating Spatial Accelerator Architectures with Tiled Matrix-Matrix Multiplication,” *IEEE Transactions on Parallel and Distributed Systems (TPDS)*, Apr. 2022.
- [3] **Geonhwa Jeong**, Eric Qin, Ananda Samajdar, Christopher J. Hughes, Sreenivas Subramoney, Hyesoon Kim, Tushar Krishna, “RSA: Efficient Register-Aware Systolic Array Matrix Engine for CPU,” in *Proceedings of the 58th Annual Design Automation Conference (DAC)*, Dec. 2021.
- [4] **Geonhwa Jeong**, Gokcen Kestor, Prasanth Chatarasi, Angshuman Parashar, Po-An Tsai, Siva Rajamanickam, Roberto Gioiosa, and Tushar Krishna, “Union: A Unified HW-SW Co-Design Ecosystem in MLIR for Evaluating Tensor Operations on Spatial Accelerators,” in *Proceedings of 30th International Conference on Parallel Architectures and Compilation Techniques (PACT)*, Sep. 2021.
- [5] Eric Qin, **Geonhwa Jeong**, William Won, Sheng-Chun Kao, Hyoukjun Kwon, Sudarshan Srinivasan, Dipankar Das, Gordon E. Moon, Sivasankaran Rajamanickam, Tushar Krishna, “Extending Sparse Tensor Accelerators to Support Multiple Compression Formats,” in *Proceedings of the 35th IEEE International Parallel Distributed Processing Symposium (IPDPS)*, May 2021.
- [6] Jan Moritz Joseph, Lennart Bamberg, **Geonhwa Jeong**, Ruei-Ting Chien, Rainer Leupers, Alberto Garcia-Ortiz, Tushar Krishna, Thilo Pionteck, “Bridging the Frequency Gap in Heterogeneous 3D SoCs through Technology-Specific NoC Router Architectures,” in *Proceedings of the 26th Asia and South Pacific Design Automation Conference (ASP-DAC)*, Jan. 2021.
- [7] Sheng-Chun Kao, **Geonhwa Jeong**, Tushar Krishna, “ConfuciuX: Autonomous Hardware Resource Assignment for DNN Accelerators using Reinforcement Learning,” in *Proceedings of 53rd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, Oct. 2020.
- [8] Kyoungmin Kim, In Seo, Wook-shin Han, Jeong-Hoon Lee, Sungpack Hong, Hassan Chafi, Hyungyu Shin, **Geonhwa Jeong**, “TurboFlux: A Fast Continuous Subgraph Matching System for Streaming Graph Data,” in *Proceedings of the 44th International Conference on Management of Data (SIGMOD)*, June 2018.

HONORS AND AWARDS

Scholarship from Kwanjeong Educational Foundation	Sep. 2019 - Current
HPCA'20 Student Travel Grant from IEEE TCCA	Feb. 2020
University Scholarship for Student Mentoring Program	Sep. 2017 - Dec. 2018
National Scholarship from ICT Creative Consilience Program	Mar. 2013 - Feb. 2019
Scholarship from Korea Hydro & Nuclear Power	Mar. 2018
Army Achievement Medal from the US Army	Aug. 2017

SKILLS

Programming	C/C++, SystemC, Java, Verilog, Ocaml
Scripting	Python, Javascript, PHP, Perl
Others	PyTorch, Tensorflow, Hadoop, Spark

PRESENTATIONS

An invited talk at Samsung Advanced Institute of Technology	May 2022
Workshop on Modeling Simulation of Systems and Applications (ModSim)	Oct. 2021
IBM AI Compute Symposium	Oct. 2021